

```
In [43]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
```

```
In [44]: from sklearn import metrics
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

Loading the data

```
In [45]: iris = pd.read_csv("iris.csv")
```

```
In [46]: print(iris.shape)
```

(150, 5)

```
In [47]: iris.head()
```

Out[47]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa


```
In [48]: iris["species"].value_counts()
```

```
Out[48]: versicolor    50  
         setosa        50  
         virginica     50  
         Name: species, dtype: int64
```

This clearly indicates that the dataset iris is clearly a balanced dataset. And each of the species contains the same number of datapoints


2DD Scatter plot

```
In [49]: iris.plot(kind='scatter', x='sepal_length', y='sepal_width') ;  
         plt.show()
```



this does not make sense out of it let us visualize more i.e use colour for each class so that we can define the things better

```
In [50]: sns.set_style("whitegrid");  
sns.FacetGrid(iris, hue="species", height=4) \  
    .map(plt.scatter, "sepal_length", "sepal_width") \  
    .add_legend();  
plt.show();
```



Notice that the blue points can be easily separated. From red and green by drawing a line. But red and green data points cannot be easily separated.

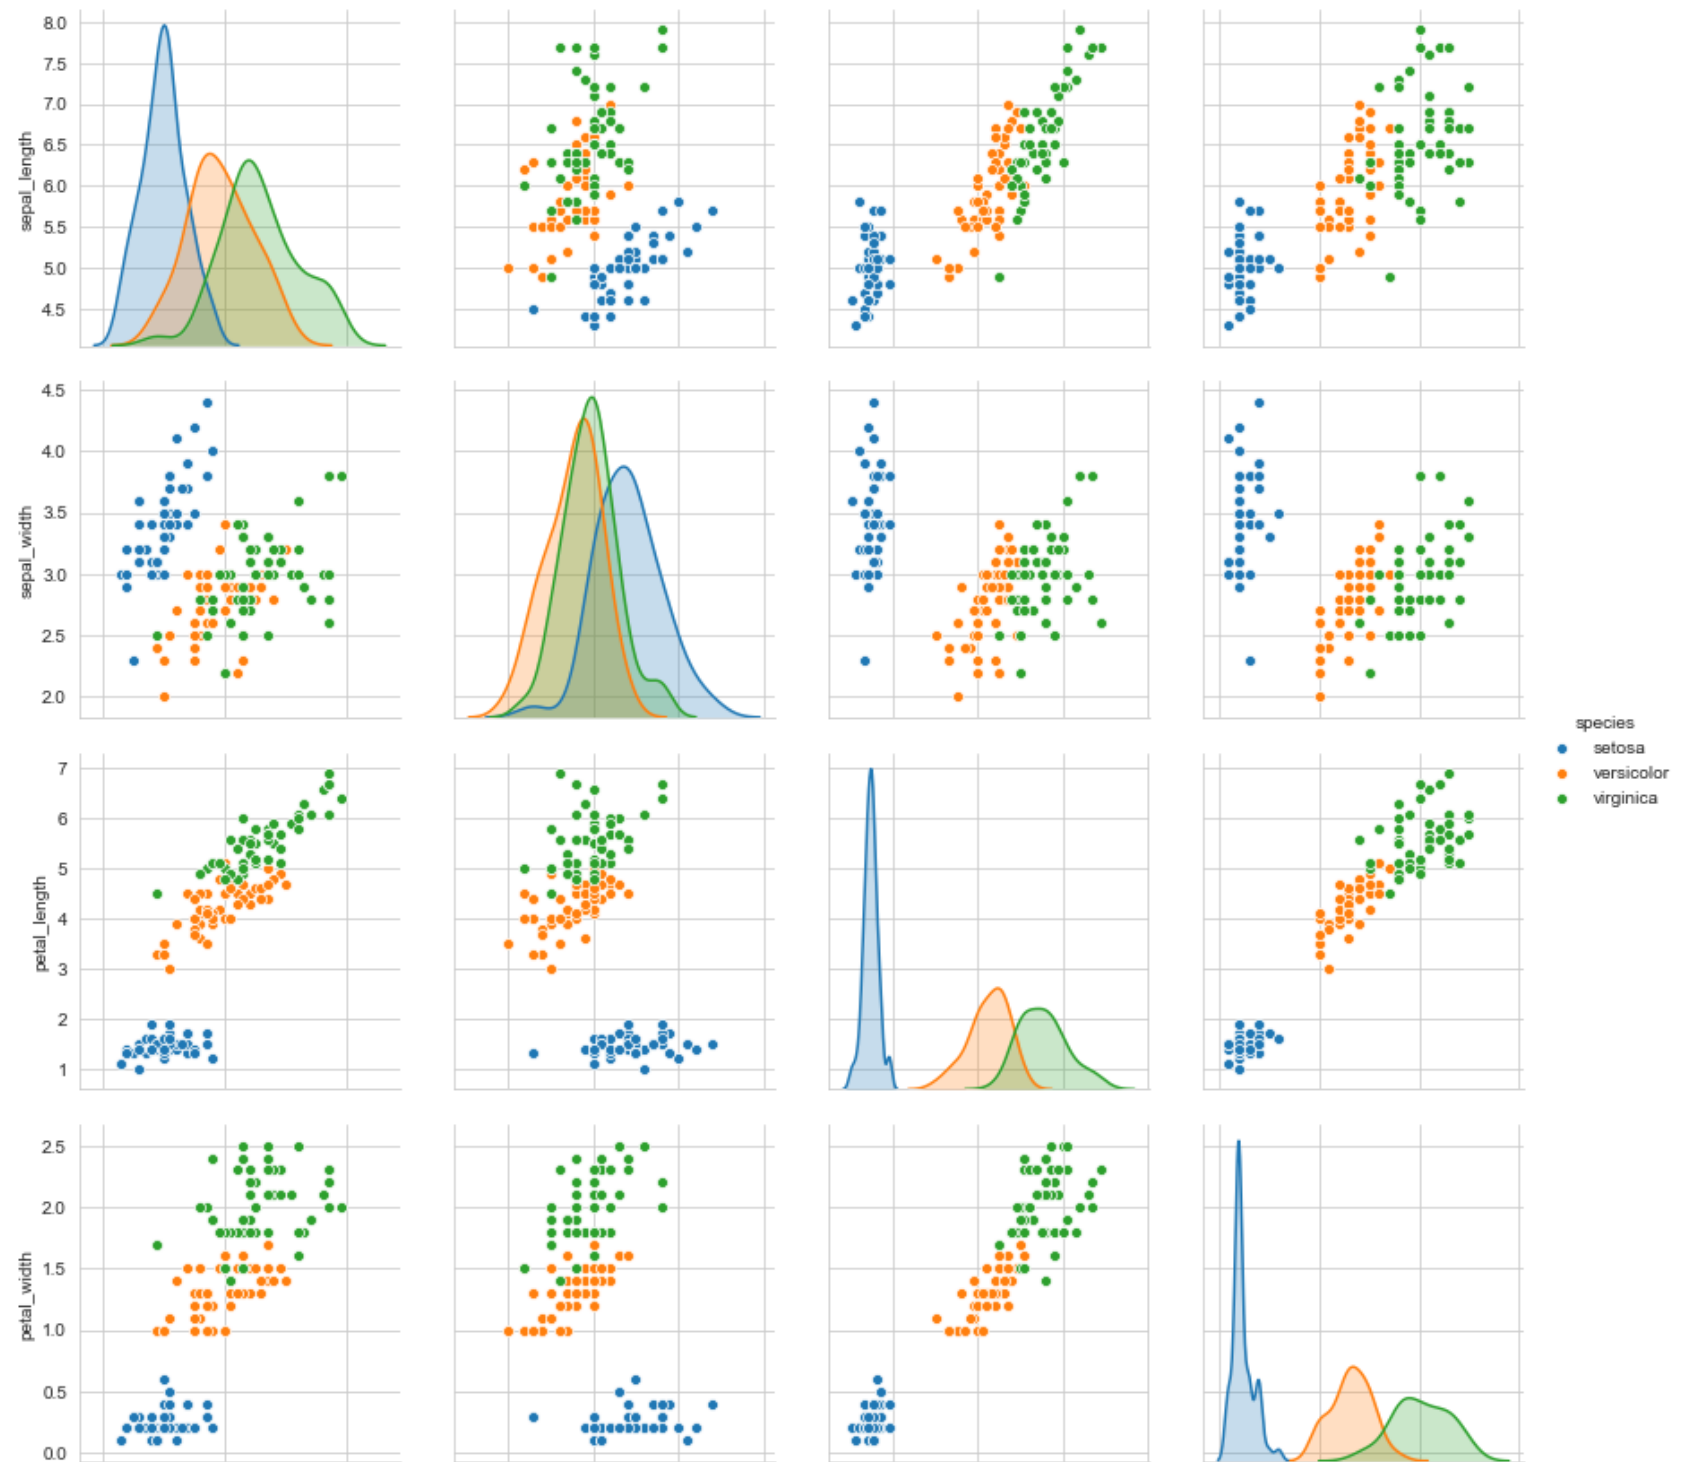
Observation


- 1: Using sepal_length and sepal_width features, we can distinguish Setosa flowers from others.
- 2: Separating Versicolor from Virginica is much harder as they have considerable overlap.

3D Scatter plot

Pair plot

```
In [51]: plt.close();  
sns.set_style("whitegrid");  
sns.pairplot(iris, hue="species", height=3);  
plt.show()
```






Observations

1. petal_length and petal_width are the most useful features to identify various flower types.
2. While Setosa can be easily identified (linearly separable), Virginica and Versicolor have some overlap (almost linearly separable).
3. We can find "lines" and "if-else" conditions to build a simple model to classify the flower types.

Histogram, PDF, CDF

```
In [52]: iris_setosa = iris.loc[iris["species"] == "setosa"];
iris_virginica = iris.loc[iris["species"] == "virginica"];
iris_versicolor = iris.loc[iris["species"] == "versicolor"];
#print(iris_setosa["petal_length"])
plt.plot(iris_setosa["petal_length"], np.zeros_like(iris_setosa['petal_length']), 'o')
plt.plot(iris_versicolor["petal_length"], np.zeros_like(iris_versicolor['petal_length']), 'o')
plt.plot(iris_virginica["petal_length"], np.zeros_like(iris_virginica['petal_length']), 'o')


plt.show()
```




Here the points are overlapping a lot.

```
In [53]: sns.FacetGrid(iris, hue="species", size=5) \
        .map(sns.distplot, "petal_length") \
        .add_legend();
plt.show();
```


C:\Users\SUNNY\Anaconda3\lib\site-packages\seaborn\axisgrid.py:230: UserWarning: The `size` paramter has been renamed to `height`; please update your code.
warnings.warn(msg, UserWarning)





```
In [54]: sns.FacetGrid(iris, hue="species", size=5) \
        .map(sns.distplot, "petal_width") \
        .add_legend();
plt.show();
```



```
In [55]: sns.FacetGrid(iris, hue="species", size=5) \
        .map(sns.distplot, "sepal_length") \
        .add_legend();
plt.show();
```



```
In [56]: sns.FacetGrid(iris, hue="species", size=5) \
        .map(sns.distplot, "sepal_width") \
        .add_legend();
plt.show();
```



```
In [57]: # Need for Cumulative Distribution Function (CDF)
# We can visually see what percentage of versicolor flowers have a


#Plot CDF of petal_length

counts, bin_edges = np.histogram(iris_setosa['petal_length'], bins=10,
                                density = True)
pdf = counts/(sum(counts))
print(pdf);
print(bin_edges);
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf);
plt.plot(bin_edges[1:], cdf)

counts, bin_edges = np.histogram(iris_setosa['petal_length'], bins=20,
                                density = True)
pdf = counts/(sum(counts))
plt.plot(bin_edges[1:],pdf);

plt.show();
```

```
[0.02 0.02 0.04 0.14 0.24 0.28 0.14 0.08 0.  0.04]
[1.   1.09 1.18 1.27 1.36 1.45 1.54 1.63 1.72 1.81 1.9 ]
```



```
In [58]: # Need for Cumulative Distribution Function (CDF)
# We can visually see what percentage of versicolor flowers have a

#Plot CDF of petal_Length


counts, bin_edges = np.histogram(iris_setosa['petal_length'], bins=10,
                                density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)

#compute CDF
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

plt.show();
```

```
[0.02 0.02 0.04 0.14 0.24 0.28 0.14 0.08 0.    0.04]
[1.    1.09 1.18 1.27 1.36 1.45 1.54 1.63 1.72 1.81 1.9 ]
```



```
In [59]: # Plots of CDF of petal_length for various types of flowers.

# Misclassification error if you use petal_length only.

counts, bin_edges = np.histogram(iris_setosa['petal_length'], bins=10,
                                density = True)
pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

# virginica
counts, bin_edges = np.histogram(iris_virginica['petal_length'], bins=10,
                                density = True)
pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)


#versicolor
counts, bin_edges = np.histogram(iris_versicolor['petal_length'], bins=10,
                                density = True)
pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:], cdf)

plt.show();
```

```

[0.02 0.02 0.04 0.14 0.24 0.28 0.14 0.08 0.  0.04]
[1.   1.09 1.18 1.27 1.36 1.45 1.54 1.63 1.72 1.81 1.9 ]
[0.02 0.1  0.24 0.08 0.18 0.16 0.1  0.04 0.02 0.06]
[4.5  4.74 4.98 5.22 5.46 5.7  5.94 6.18 6.42 6.66 6.9 ]
[0.02 0.04 0.06 0.04 0.16 0.14 0.12 0.2  0.14 0.08]
[3.   3.21 3.42 3.63 3.84 4.05 4.26 4.47 4.68 4.89 5.1 ]

```



Mean, Variance and Std-dev

```
In [60]: #Mean, Variance, Std-deviation,  
print("Means:")  
print(np.mean(iris_setosa["petal_length"]))  
#Mean with an outlier.  
print(np.mean(np.append(iris_setosa["petal_length"],50)));  
print(np.mean(iris_virginica["petal_length"]))  
print(np.mean(iris_versicolor["petal_length"]))  
  
print("\nStd-dev:");  
print(np.std(iris_setosa["petal_length"]))  
print(np.std(iris_virginica["petal_length"]))  
print(np.std(iris_versicolor["petal_length"]))
```

Means:

1.464
2.4156862745098038
5.552
4.26

Std-dev:

0.17176728442867115
0.5463478745268441
0.4651881339845204

Median, Percentile, Quantile, IQR, MAD


```
In [61]: #Median, Quantiles, Percentiles, IQR.
print("\nMedians:")
print(np.median(iris_setosa["petal_length"]))
#Median with an outlier
print(np.median(np.append(iris_setosa["petal_length"],50)));
print(np.median(iris_virginica["petal_length"]))
print(np.median(iris_versicolor["petal_length"]))

print("\nQuantiles:")
print(np.percentile(iris_setosa["petal_length"],np.arange(0, 100, 25)))
print(np.percentile(iris_virginica["petal_length"],np.arange(0, 100, 25)))
print(np.percentile(iris_versicolor["petal_length"], np.arange(0, 100, 25)))

print("\n90th Percentiles:")
print(np.percentile(iris_setosa["petal_length"],90))
print(np.percentile(iris_virginica["petal_length"],90))
print(np.percentile(iris_versicolor["petal_length"], 90))

from statsmodels import robust
print ("\nMedian Absolute Deviation")
print(robust.mad(iris_setosa["petal_length"]))
print(robust.mad(iris_virginica["petal_length"]))
print(robust.mad(iris_versicolor["petal_length"]))
```

Medians:

1.5

1.5

5.55

4.35

Quantiles:

[1. 1.4 1.5 1.575]

[4.5 5.1 5.55 5.875]

[3. 4. 4.35 4.6]

90th Percentiles:

1.7

6.3100000000000005

4.8

Median Absolute Deviation

0.14826022185056031

0.6671709983275211

0.5189107764769602


Box plot and Whiskers

```
In [62]: #Box-plot with whiskers: another method of visualizing the 1-D scatter plot more intuitively.  
# The Concept of median, percentile, quantile.
```

```
# IN the plot below, a technique call inter-quartile range is used in plotting the whiskers.  
#Whiskers in the plot below donot correposnd to the min and max values.
```


```
#Box-plot can be visualized as a PDF on the side-ways.
```

```
sns.boxplot(x='species',y='petal_length', data=iris)  
plt.show()
```




violin plots

```
In [63]: # A violin plot combines the benefits of the previous two plots  
#and simplifies them  
  
# Denser regions of the data are fatter, and sparser ones thinner  
#in a violin plot  
  
sns.violinplot(x="species", y="petal_length", data=iris, size=8)  
plt.show()
```



```
In [64]: #2D Density plot, contours-plot  
sns.jointplot(x="petal_length", y="petal_width", data=iris_setosa, kind="kde");  
plt.show();
```



Modeling with scikit learn

```
In [76]: X = iris.drop([ 'species'], axis=1)
y = iris['species']
print(X.head())
print("_"*100)

print(X.shape)
print("_"*100)
print(y.head())
print("_"*100)
print(y.shape)
```

	sepal_length	sepal_width	petal_length	petal_width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

(150, 4)

0	setosa
1	setosa
2	setosa
3	setosa
4	setosa


Name: species, dtype: object

(150,)

Initialization the knn model

```
In [66]: k_range = list(range(1,26))
scores = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X, y)
    y_pred = knn.predict(X)
    scores.append(metrics.accuracy_score(y, y_pred))

plt.plot(k_range, scores)
plt.xlabel('Value of k for KNN')
plt.ylabel('Accuracy Score')
plt.title('Accuracy Scores for Values of k of k-Nearest-Neighbors')
plt.show()
```



Initialization Logistic Regression

```
In [67]: logreg = LogisticRegression()  
logreg.fit(X, y)  
y_pred = logreg.predict(X)  
print(metrics.accuracy_score(y, y_pred))
```

0.96

C:\Users\SUNNY\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\SUNNY\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:469: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)

By splitting the dataset pseudo-randomly into a two separate sets, we can train using one set and test using another.

This ensures that we won't use the same observations in both sets. More flexible and faster than creating a model using all of the dataset for training.

Spiting the data into train and test data

```
In [68]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=5)  
print(X_train.shape)  
print(y_train.shape)  
print(X_test.shape)  
print(y_test.shape)
```


(105, 4)

(105,)

(45, 4)

(45,)


```
In [77]: ##  
         #initialization the KNN MODEL  
  
         k_range = list(range(1,26))  
         scores = []  
         for k in k_range:  
             knn = KNeighborsClassifier(n_neighbors=k)  
             knn.fit(X_train, y_train)  
             y_pred = knn.predict(X_test)  
             scores.append(metrics.accuracy_score(y_test, y_pred))  
  
         plt.plot(k_range, scores)  
         plt.xlabel('Value of k for KNN')  
         plt.ylabel('Accuracy Score')  
         plt.title('Accuracy Scores for Values of k of k-Nearest-Neighbors')  
         plt.show()
```



In [70]: *## initialization Logistic regression*

```
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)  
y_pred = logreg.predict(X_test)  
print(metrics.accuracy_score(y_test, y_pred))
```

0.9333333333333333

C:\Users\SUNNY\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\SUNNY\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:469: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)

In []:

In []:

In []:

In []: